

Generative AI - Text to Text - Part 2



Ahmed Elbagoury

Introduction – Ahmed Elbagoury

Senior ML Engineer, Google.



ML Research + Teaching



Applied ML



Teaching Applied ML

Optimize Your Experience

- ✓ Ask questions in the class
- ✓ Don't ignore the math (Whys and Hows)
- ✓ Fortify understanding by reading papers, walking through code implementations and trying out notebooks!



- ✓ Enjoy the subject. Study deeply, seek understanding, read papers, practice problems, ask questions
- ✓ Don't spend time in self-doubt. Unlike real life tortoise & hare race, slow & steady literally wins the career race

Here's the plan

Part 1



How LLM can be used in applications



Part 2



Decoder details and Optimizations

⇐ **Today**

Today's Agenda



1



GPT Models

2



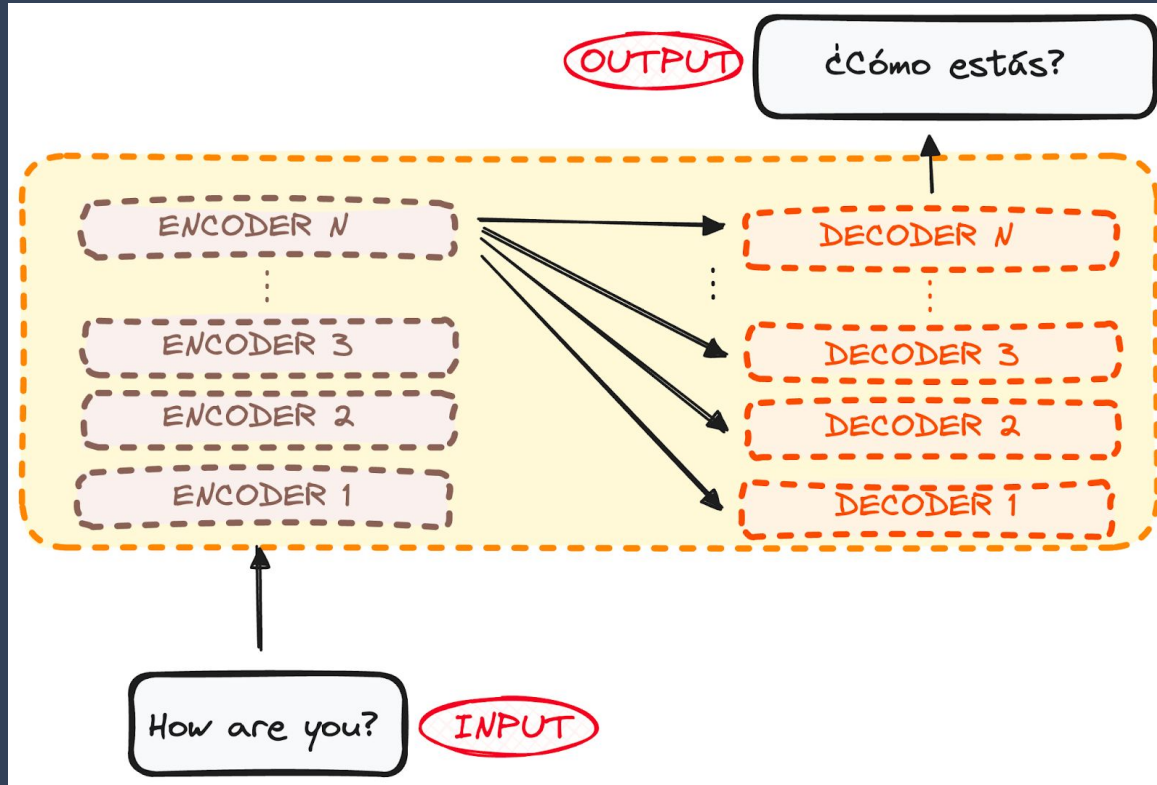
Sampling
Strategies

3



LLM
Optimizations

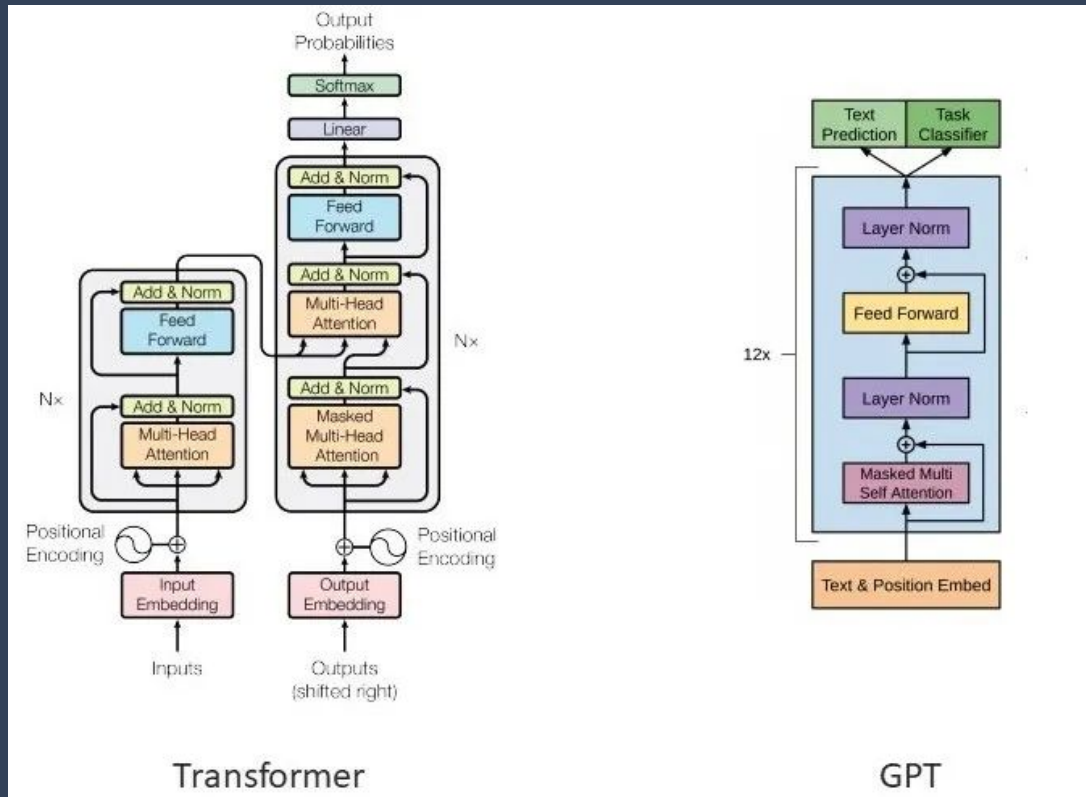
We talked about encoder-decoder models



The diagram illustrates the BERT architecture. At the bottom, the input tokens are listed: [CLS], Help, Prince, Mayuko, followed by an ellipsis and 512. Above these tokens are three stacked encoder blocks, labeled 1, 2, and 12 from bottom to top. Each encoder block is a green rounded rectangle containing the word 'ENCODER'. A yellow rounded rectangle encloses the entire stack of encoder blocks. Arrows point from each token to its corresponding encoder block. Above the encoder stack, there are five light blue boxes representing the output of the encoder stack. The first box, corresponding to the [CLS] token, is highlighted in blue. The other four boxes are light blue. The tokens 'Help', 'Prince', and 'Mayuko' are highlighted in green in the input list.

What is missing?

Decoder Models





Quiz

Which Component should be dropped from the decoder in GPT (decoder only) models?


Quiz

Which Component should be dropped from the decoder in GPT (decoder only) models?

⇒ Encoder-decoder attention should be dropped since there is no encoder!

Decoder Models - Text Prediction

“Not all heroes wear capes” \Rightarrow 5

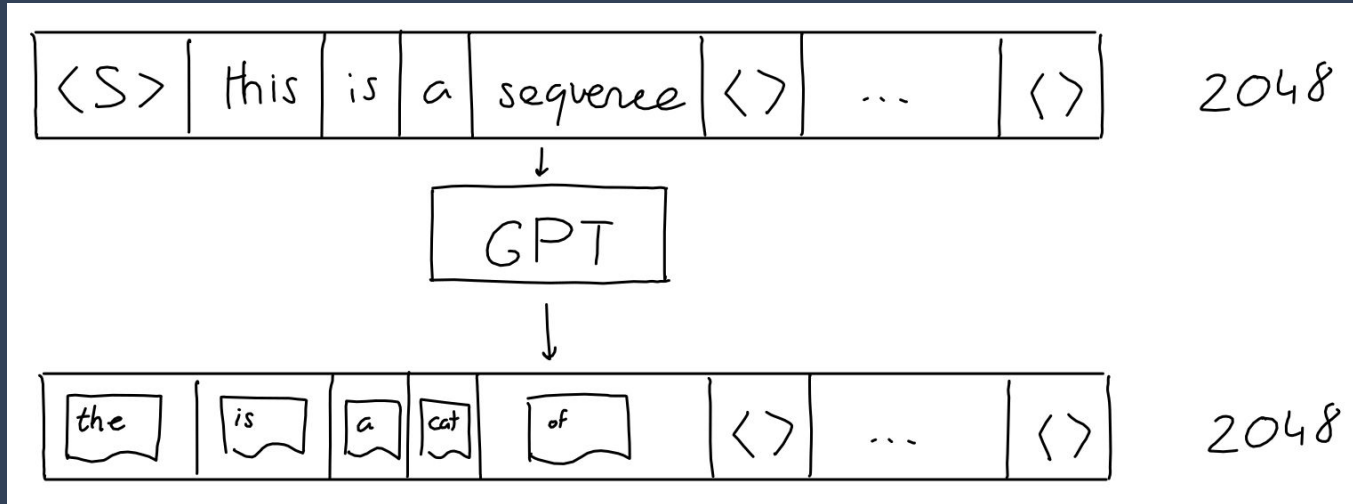


Decoder Models - Text Prediction

How to get multiple words and for different lengths?!

Decoder Models - Text Prediction



How to get multiple words and for different lengths?!



When generating text, we typically only look at the guess for the last word of the sequence.

Decoder Models - Task Classifier

After pre-training


- GPT can capture linguistic knowledge of input sequences 
- However, to make it better perform on downstream tasks 

⇒ it needs to be fine-tuned on a supervised problem.

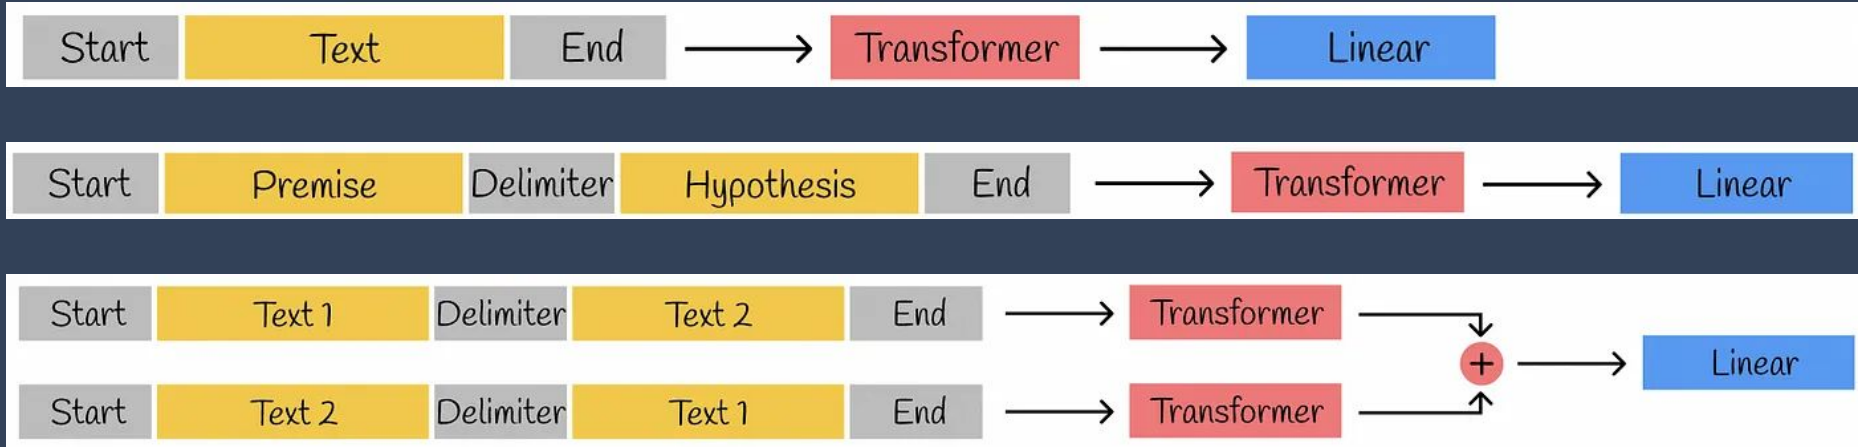
Decoder Models - Task Classifier





Decoder Models - Task Classifier



Decoder Models - Task Classifier



Decoder Models - Task Classifier



Why Decoder Models aka Causal Decoder/Models?

- Cost of training for Causal Decoder (CD) is cheaper
- CD works better for In-Context learning
 - ⇒ It has a more straightforward effect for CD

**Since Decoding is very
important
Let's zoom in to it**



Today's Agenda

1



GPT Models

2




Sampling
Strategies

3

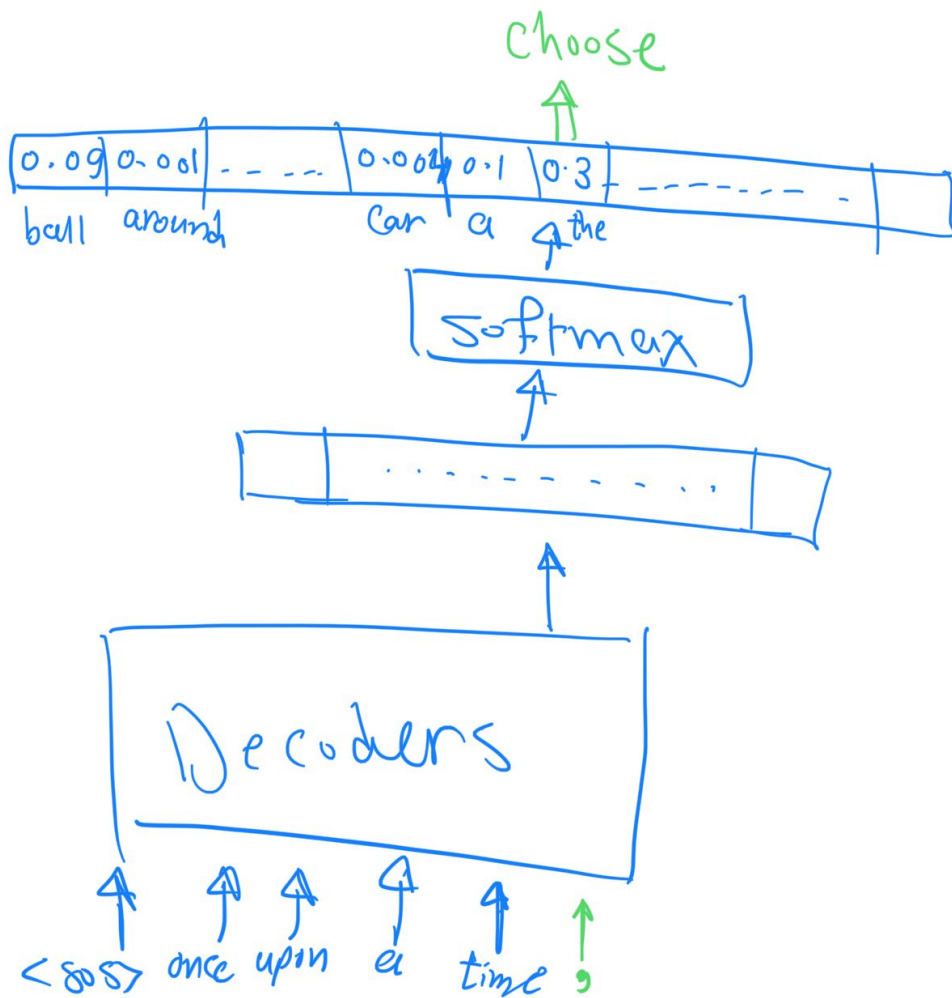


LLM
Optimizations

Decoding Details



Decoding Details



What is the Optimal Solution

- We want to maximize the likelihood of the output sequence
 - This is the multiplication of conditional probabilities
 - Or minimizing the summation of log likelihood

What is the Optimal Solution

- We want to maximize the likelihood of the output sequence
 - This is the multiplication of conditional probabilities
 - Or minimizing the summation of log likelihood

What are the problems with the previous approach? (greedy)

total = 4.2

I am a $\xrightarrow{0.3}$ student $\xrightarrow{1.9} \cdot \xrightarrow{2} \text{EOS}$

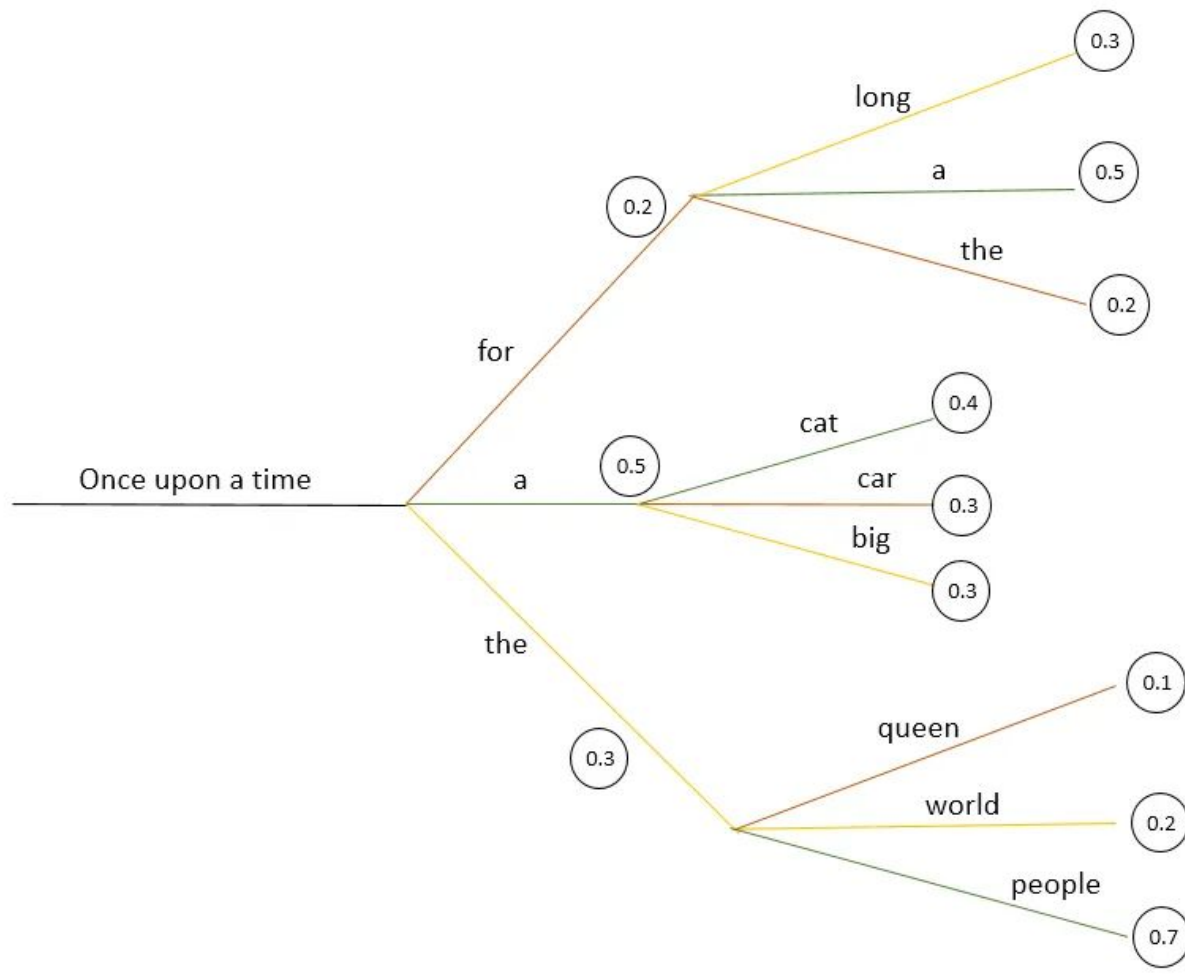
total = 1.2

$\xrightarrow{0.7}$ player $\xrightarrow{0.2}$ eats $\xrightarrow{0.3} \text{EOS}$
 $\xrightarrow{0.1}$ xyz
 $\xrightarrow{0.2}$ abc

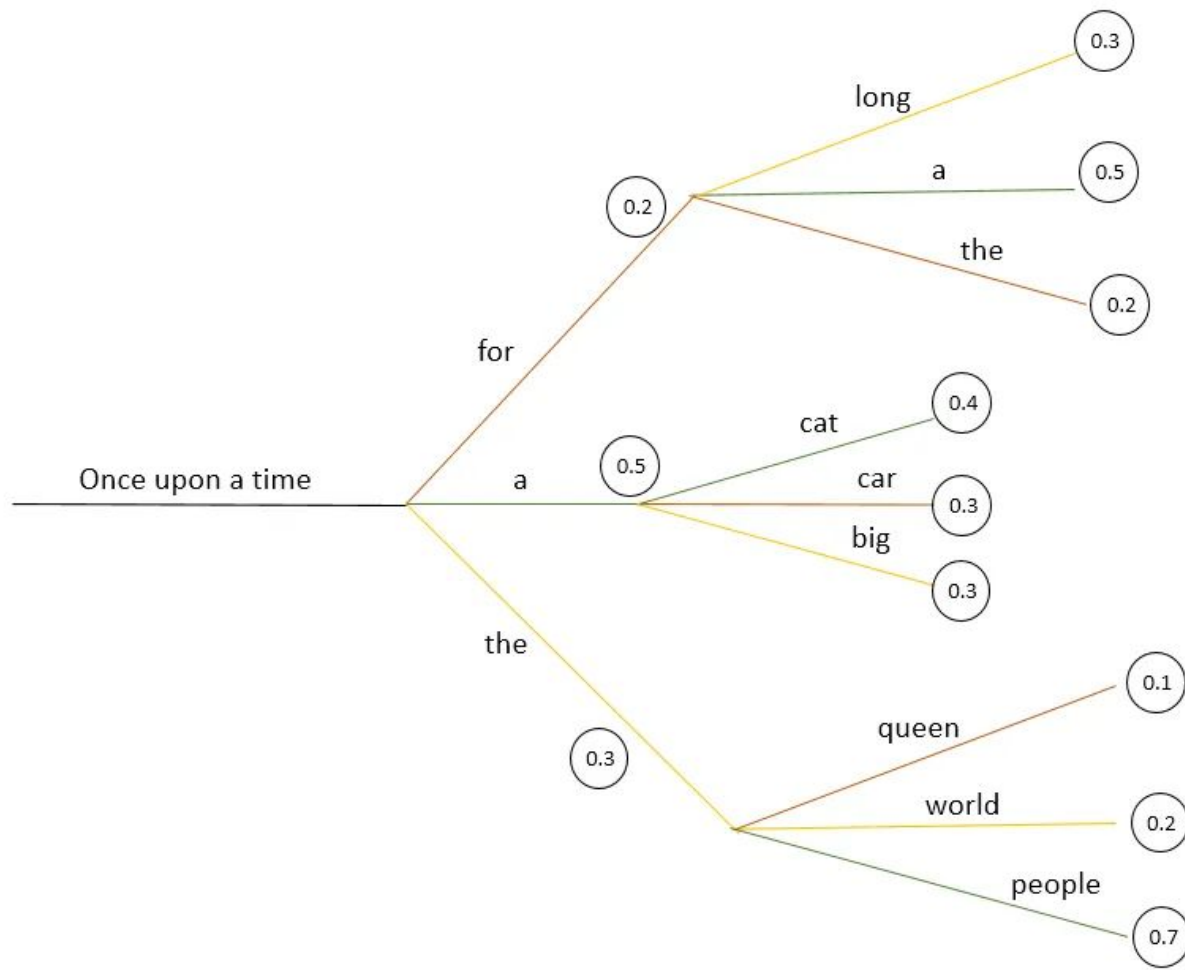
Decoding Details: Beam Search

- Keep your options open (to some extent)
- Just in case your local choices are not optimal
- It requires more memory and computational power to keep track of the beam

beam = ??



beam = 3

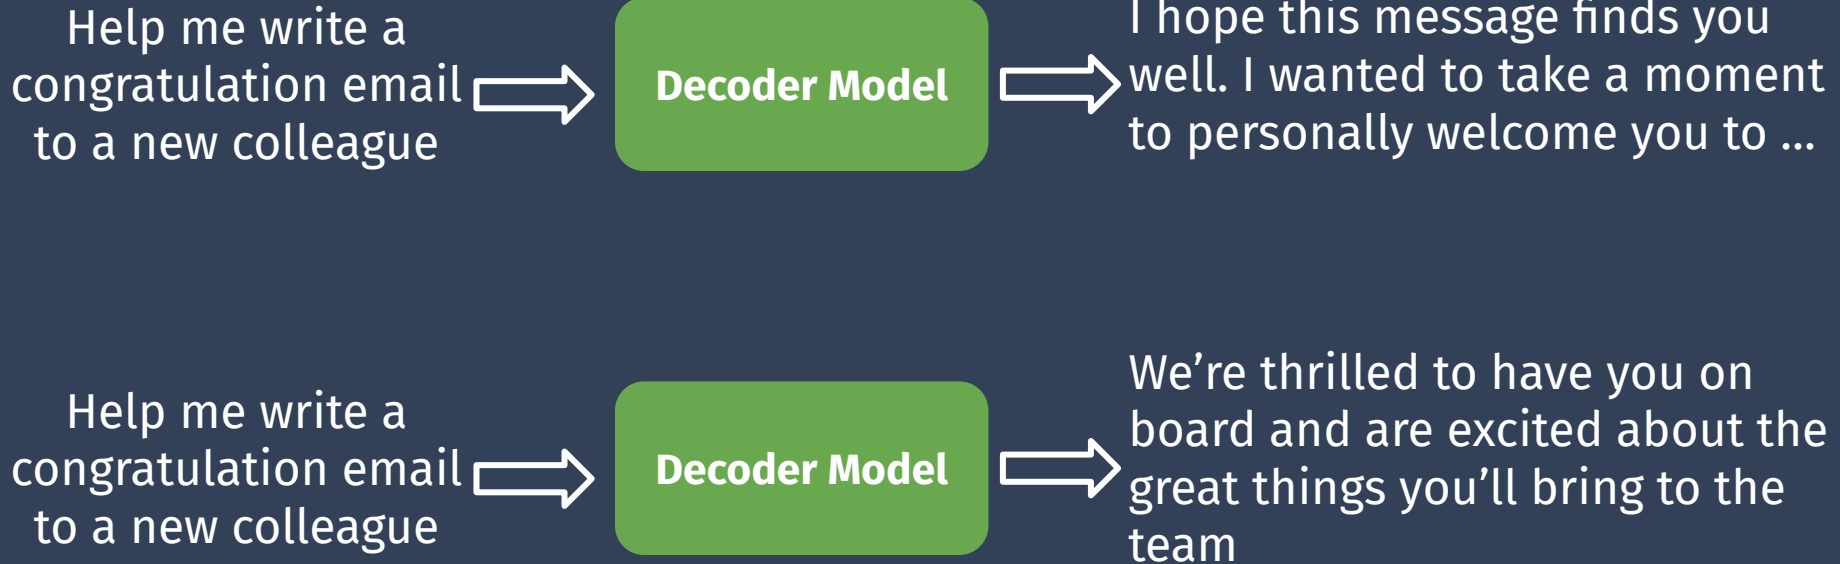


Decoding Details: What if we got the answer wrong



Let's say I want to see other drafts!

Decoding Details: What if we got the answer wrong



Decoding Details: What if we got the answer wrong

- Random (aka stochastic) sampling to the rescue

**Sorry more terminology!
So that you know the lingo!**



Decoding Details: What if we got the answer wrong

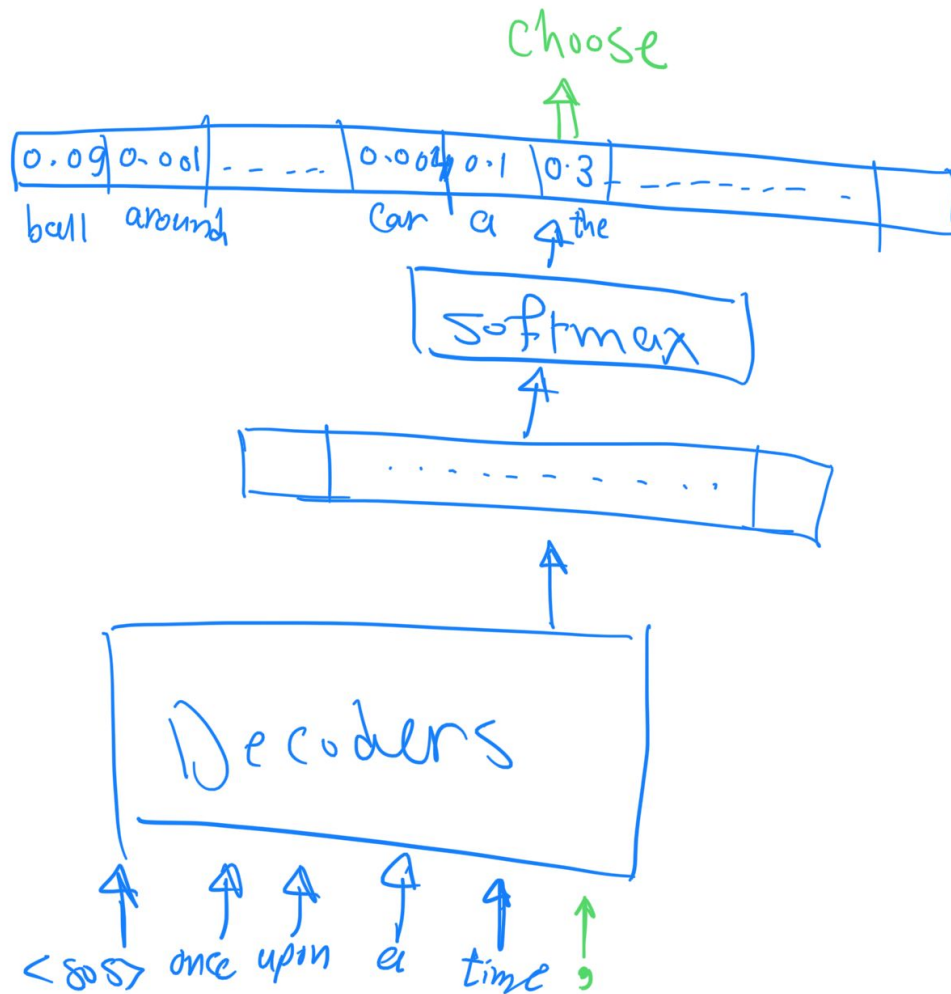
- Random (aka stochastic) sampling to the rescue

**Sorry more terminology!
So that you know the lingo!**

- It just means
 - Choose randomly according to the probability scores

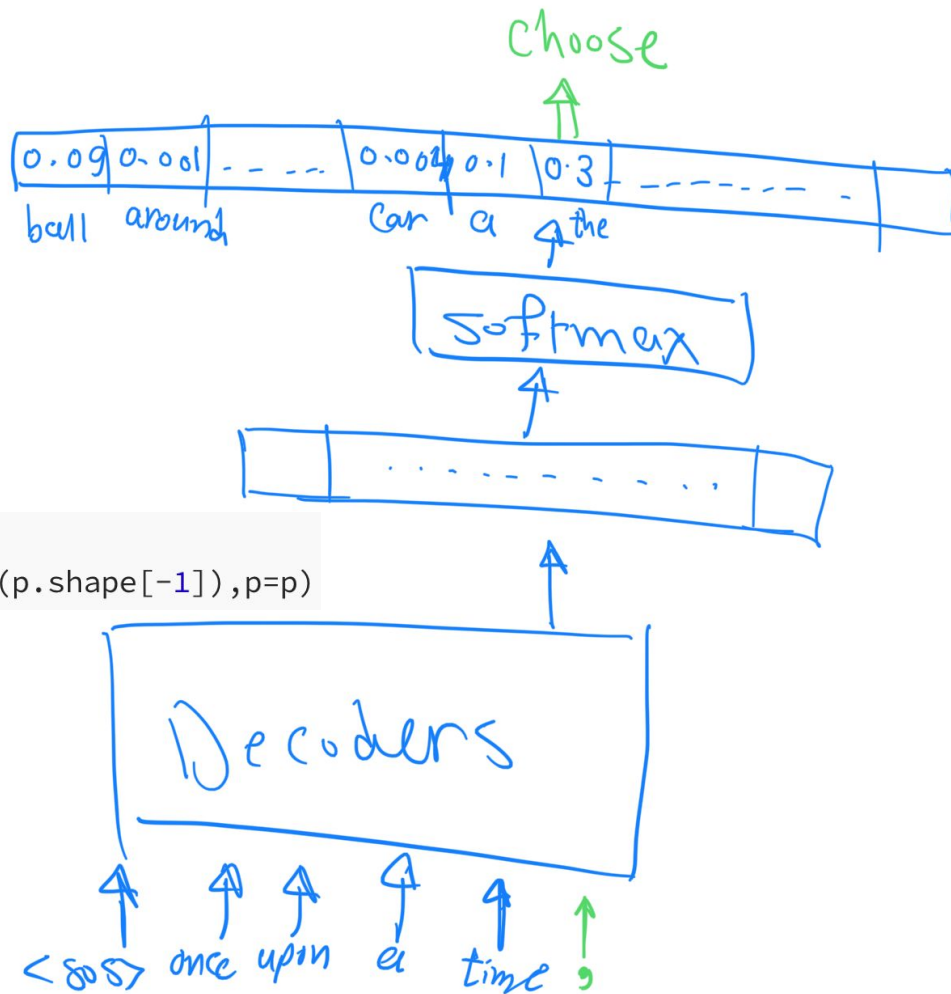


Random Sampling




Random Sampling


```
def sample(p):  
    return np.random.choice(np.arange(p.shape[-1]), p=p)
```




Decoding Details: Top-K



sample from these
potentially with
temperature




Decoding Details: Top-p aka Nucleus



Temperature

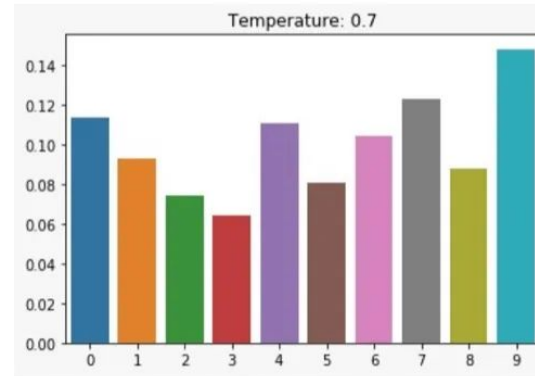
SOFTMAX WITHOUT TEMPERATURE ($T=1$)

$$\frac{e^{z_i}}{\sum_j e^{z_j}}$$



SOFTMAX WITH TEMPERATURE

$$\frac{e^{z_i/T}}{\sum_j e^{z_j/T}}$$



LESS ENTROPY

INCREASE IN ENTROPY
WITH INCREASE IN T

MORE ENTROPY

Demo Time

<https://colab.research.google.com/drive/1CejK4UaV0O6L3c4jnKqtWENp9YFboNlg>

Today's Agenda

1



GPT Models

2



Sampling
Strategies

3



LLM
Optimizations



Optimization Strategies for LLM

Larger Models Mean More Capacity

- To handle variety of tasks \Rightarrow Models need to be trained on large amount of data
- To avoid underfitting \Rightarrow Models need to be larger!

Number of Parameters

(in Millions)



Larger Models Mean More Capacity

- To handle variety of tasks
 - Models need to be trained on large amount of data
- To avoid underfitting
 - Models need to be larger!

Implications of Increasing Model Size



Implications of Increasing Model Size

- More hardware resources
 - On-device use cases?
- Higher Latency
- Worse Carbon footprint

Do we have to use these huge models?

- A line of research motivated by the challenges of training and productionizing LLMs
- It aims to
 - Improving training and/or inference latency
 - Reducing model sizes with little to no impact on quality

Do we have to use these huge models?

- A line of research motivated by the challenges of training and productionizing LLMs
- It aims to
 - Improving training and/or inference latency
 - Reducing model sizes with little to no impact on quality

Major techniques:

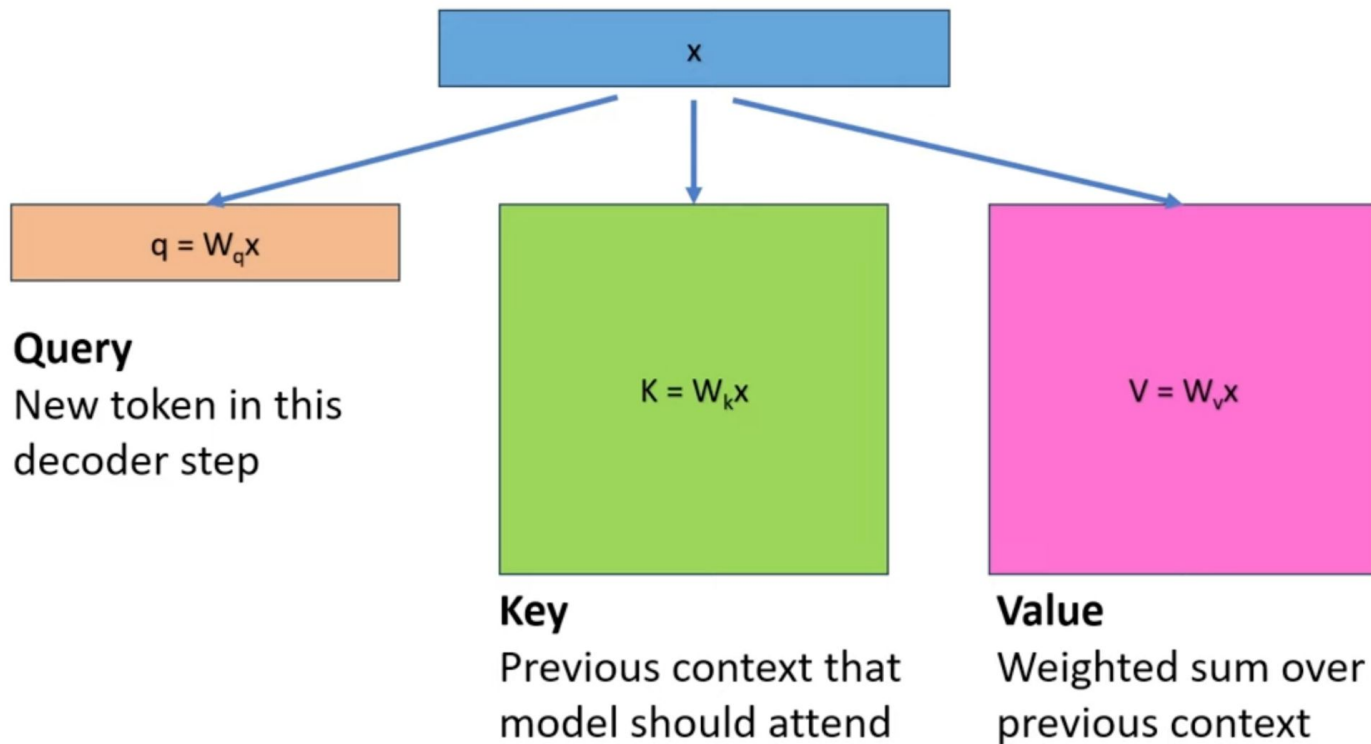
- KV Caching
- LoRA Tuning (More generally PEFT)
- Distillation
- Model Pruning
- Quantization

KV Caching

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

KV Caching

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



KV Caching

It was a cold windy morning when I stepped outside, feeling a chill

$$q = W_q x$$

$$K = W_k x$$

$$V = W_v x$$

KV Caching

It was a cold windy morning when I stepped outside, feeling a chill


$$q = W_q x$$

$$K = W_k x$$

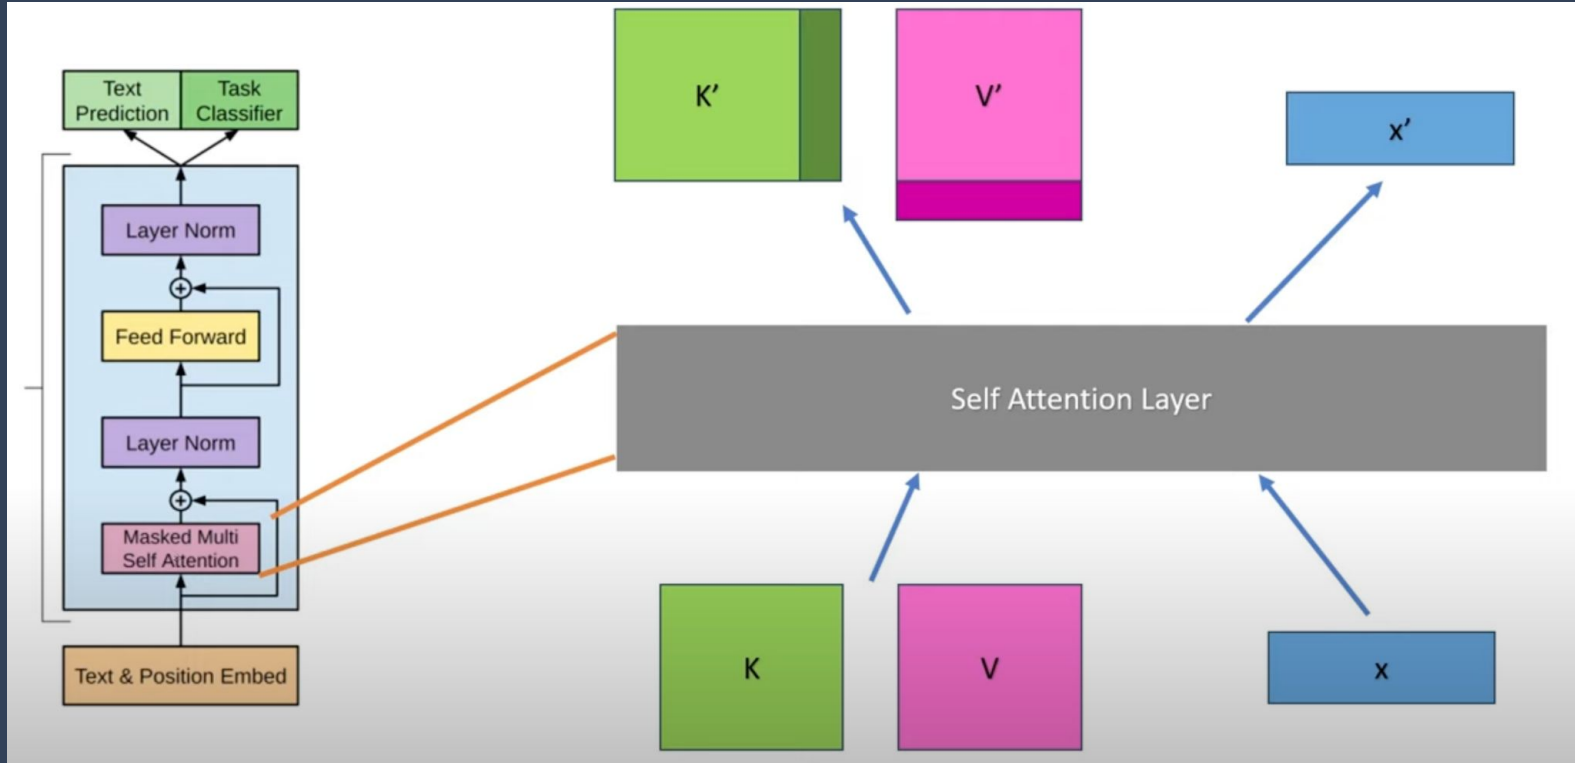
$$V = W_v x$$

KV Caching

It was a cold windy morning when I stepped outside, feeling a chill



KV Caching



Demo Time

<https://colab.research.google.com/drive/12ioUtylE5BuWTNjDHdecNQ8Xnb1dRISW>



Quiz

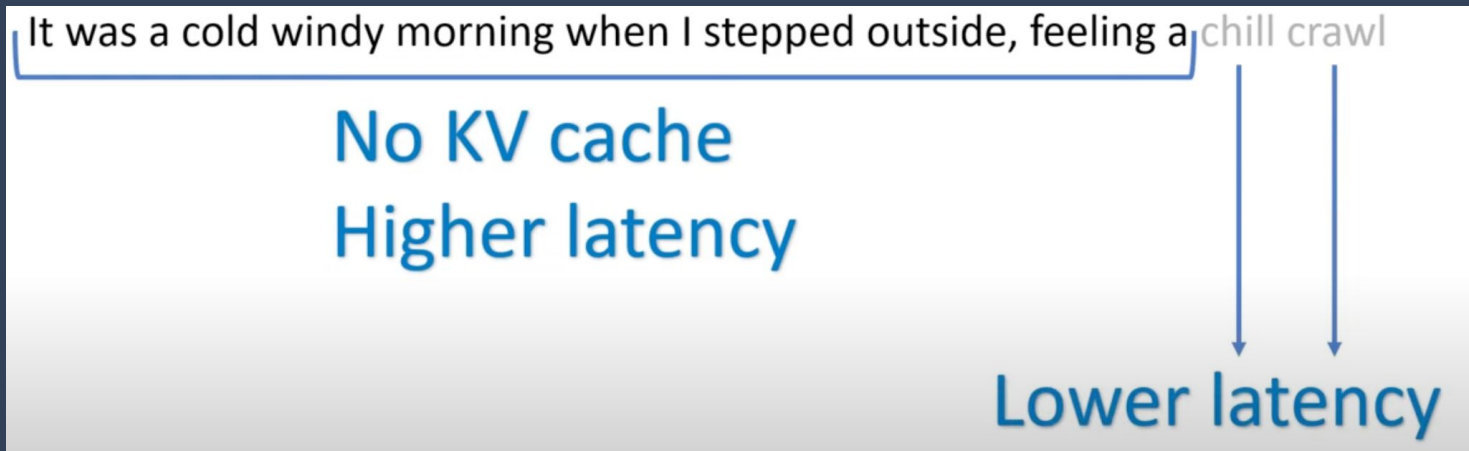
Which part will run faster: Red or Blue? and why?
Assuming that red is the prompt

It was a cold windy morning when I stepped outside, feeling a chill crawl

Quiz

Which part will run faster: Red or Blue? and why?
Assuming that red is the prompt

It was a cold windy morning when I stepped outside, feeling a chill crawl



One more Quiz

KV Caching improves inference latency at the expense of

- A) Memory Usage during inference
- B) Training Time
- C) A & B
- D) Memory Usage during training

Quiz


KV Caching improves inference latency at the expense of

- A) Memory Usage during inference
- B) Training Time
- C) A & B
- D) Memory Usage during training

LoRA

- Low-Rank Adaptation of Large Language Models
- Is a type of what is known as Parameter-Efficient tuning (PEFT)
- The goal is fine-tuning a model
 - By updating smaller number of parameters

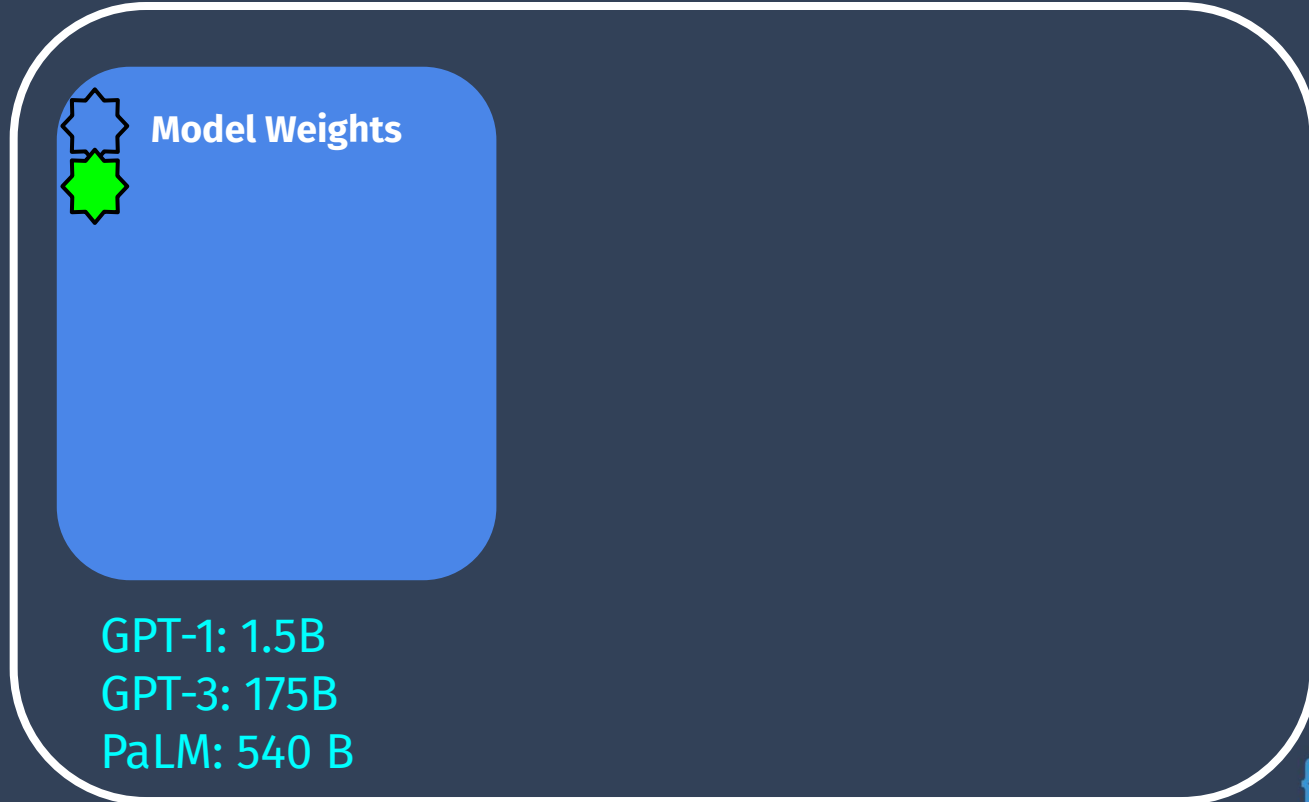
Why LoRA?



GPT-1: 1.5B
GPT-3: 175B
PaLM: 540 B


Regular Fine-Tuning

During fine-tuning we need to load this in Memory



Regular Fine-Tuning

During fine-tuning we need to load this in Memory \Rightarrow



Regular Fine-Tuning

During fine-tuning we need to load this in Memory \Rightarrow



Model Weights

$$\begin{pmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{pmatrix}$$


Gradient for each Weight

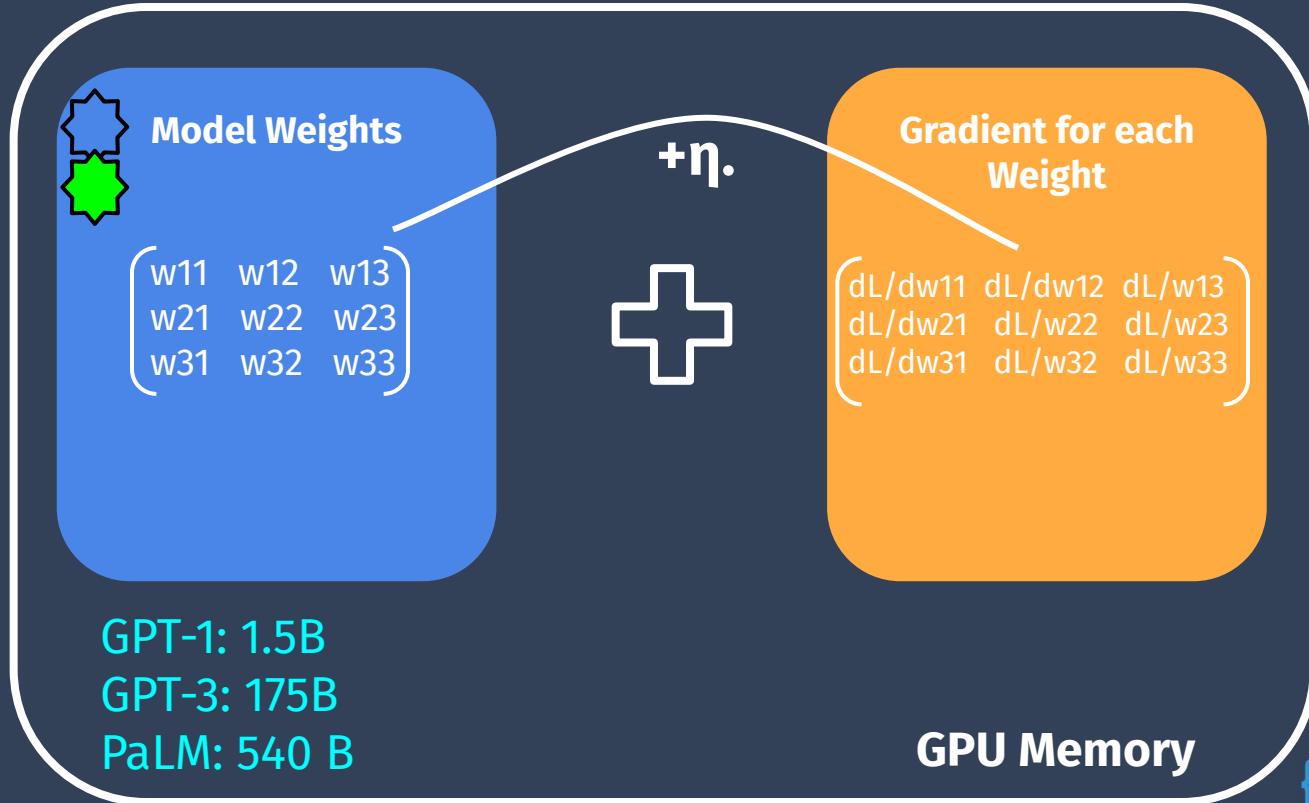
$$\begin{pmatrix} dL/dw_{11} & dL/dw_{12} & dL/w_{13} \\ dL/dw_{21} & dL/w_{22} & dL/w_{23} \\ dL/dw_{31} & dL/w_{32} & dL/w_{33} \end{pmatrix}$$

GPT-1: 1.5B
GPT-3: 175B
PaLM: 540 B

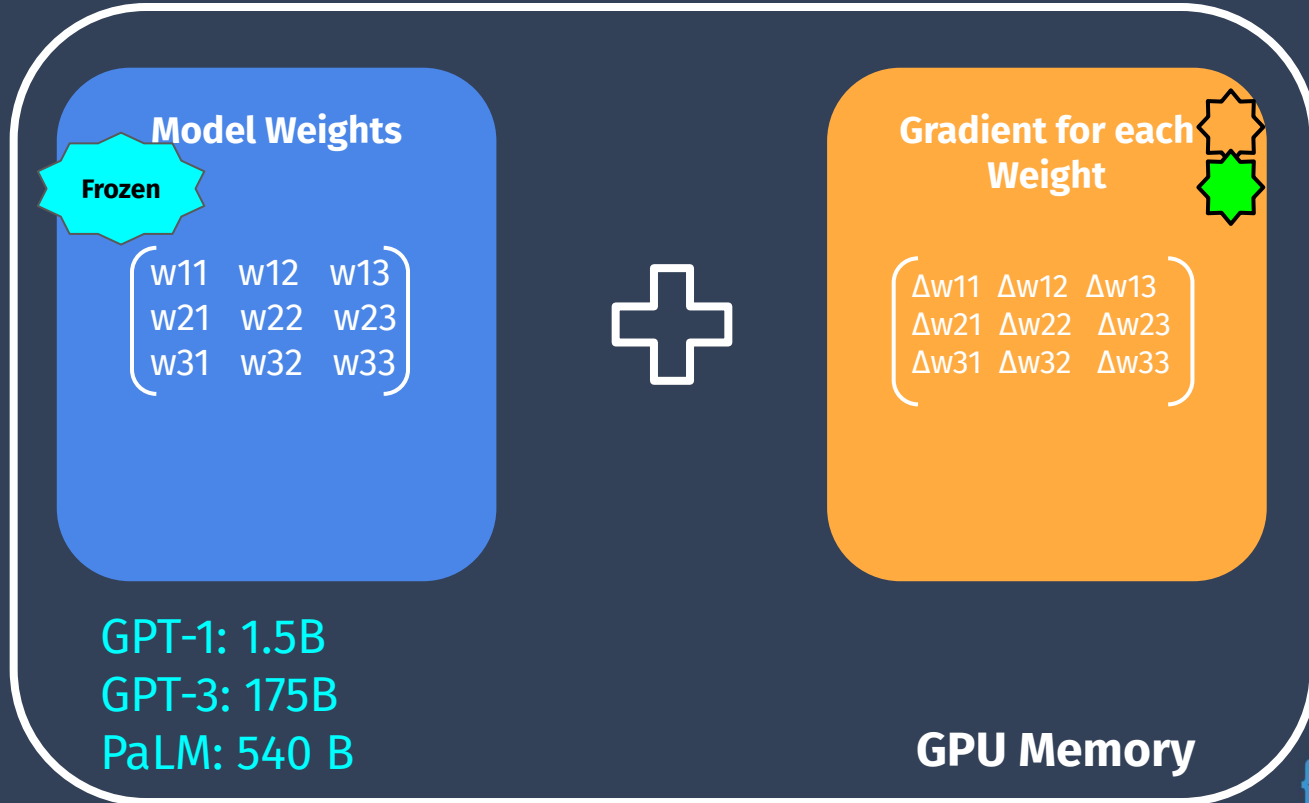
GPU Memory

Regular Fine-Tuning

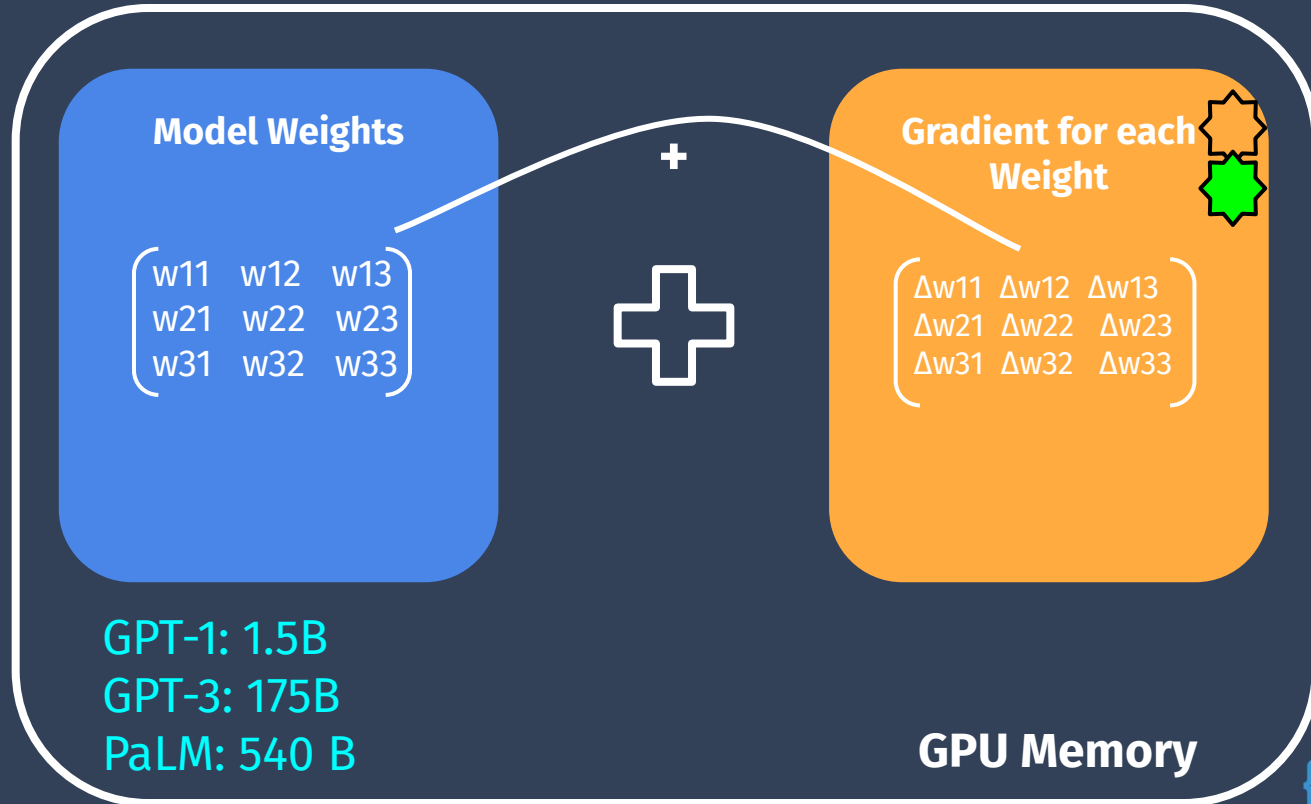
During fine-tuning we need to load \Rightarrow



LoRA

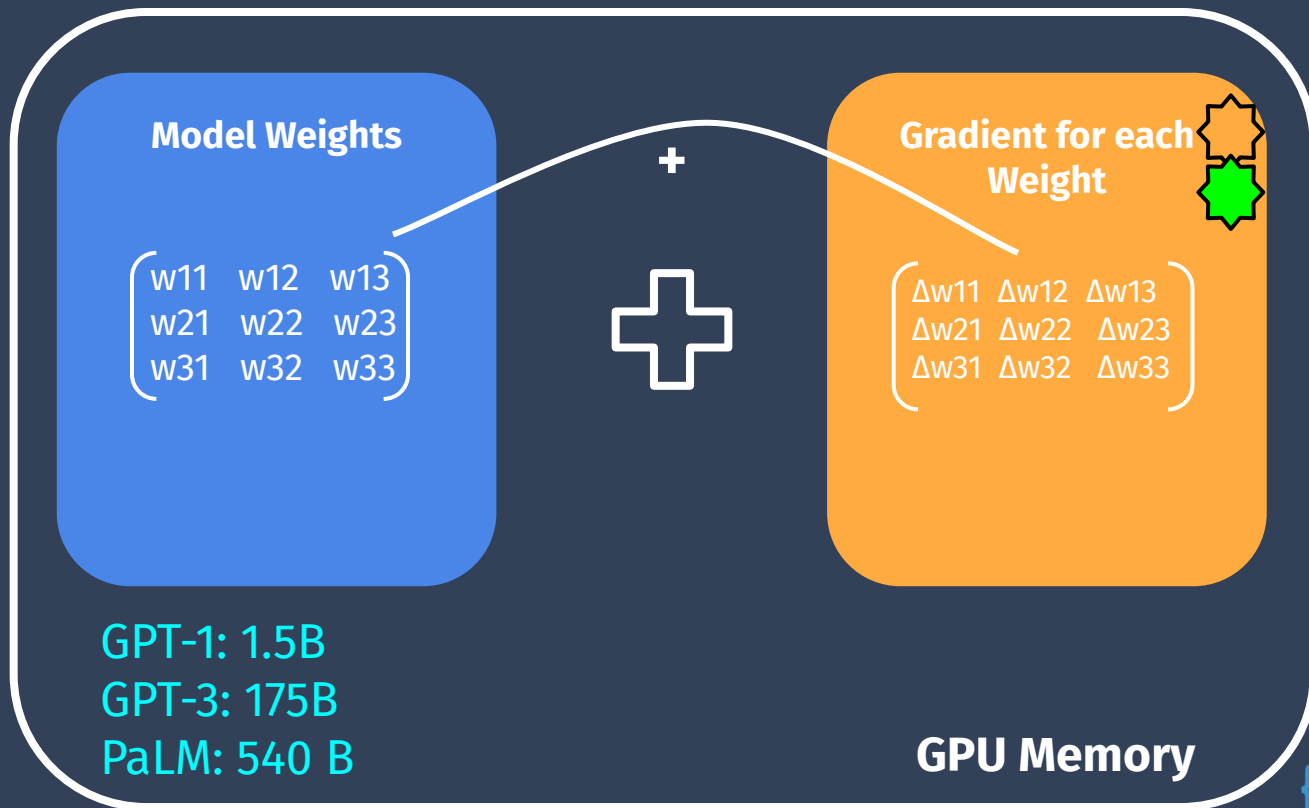


LoRA



LoRA

Does not it this double the number of parameters as well??



Does not it this double the
number of parameters as well??
No, not really ... at inference:

LoRA

Model Weights

$$\begin{pmatrix} w_{11}+\Delta w_{11} & w_{12}+\Delta w_{12} & w_{13}+\Delta w_{13} \\ w_{21}+\Delta w_{21} & w_{22}+\Delta w_{22} & w_{23}+\Delta w_{23} \\ w_{31}+\Delta w_{31} & w_{32}+\Delta w_{32} & w_{33}+\Delta w_{33} \end{pmatrix}$$

=

Model Weights

$$\begin{pmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{pmatrix}$$

+

Gradient for each Weight

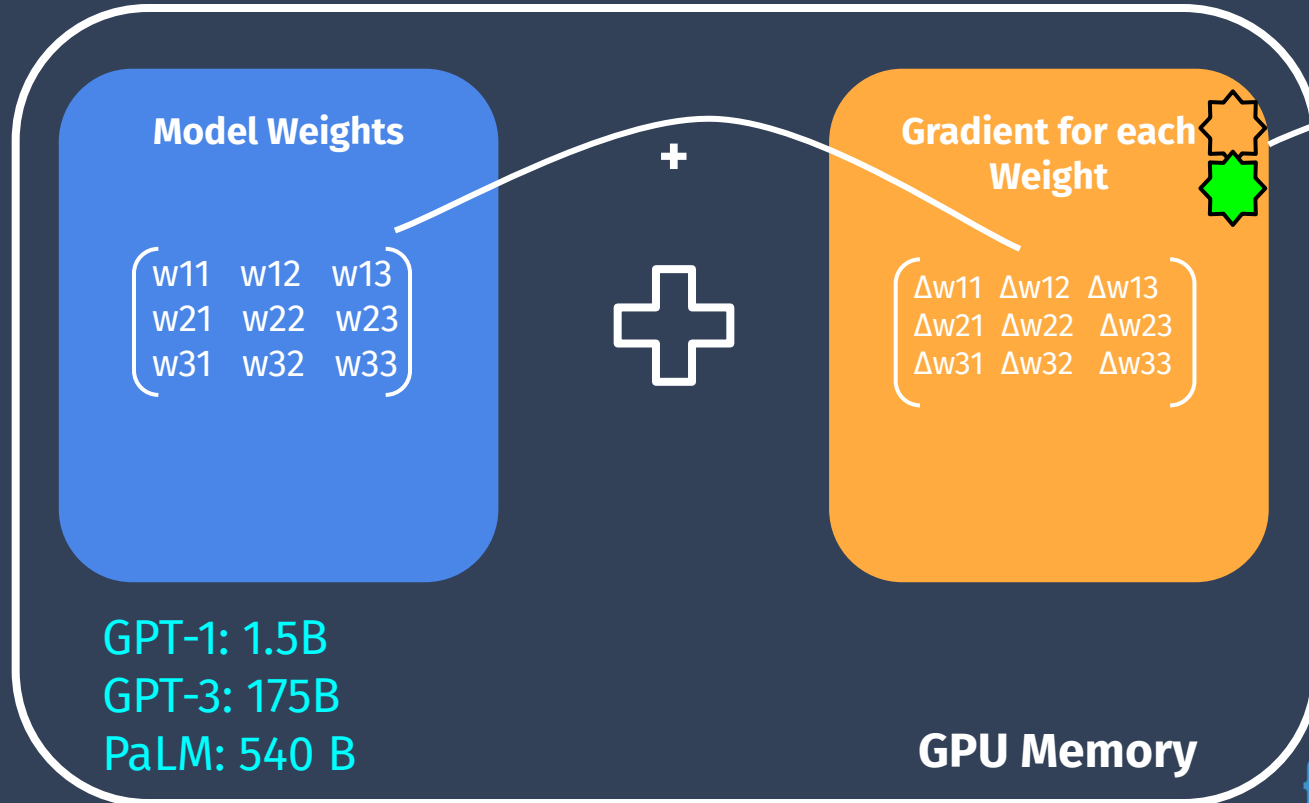
$$\begin{pmatrix} \Delta w_{11} & \Delta w_{12} & \Delta w_{13} \\ \Delta w_{21} & \Delta w_{22} & \Delta w_{23} \\ \Delta w_{31} & \Delta w_{32} & \Delta w_{33} \end{pmatrix}$$

GPT-1: 1.5B
GPT-3: 175B
PaLM: 540 B

GPU Memory

Ok but we still need to tune that part?
At training:

LoRA



LoRA Tuning

$$\Delta W = \begin{bmatrix} \Delta w_{11} & \Delta w_{12} & \Delta w_{13} \\ \Delta w_{21} & \Delta w_{22} & \Delta w_{23} \\ \Delta w_{31} & \Delta w_{32} & \Delta w_{33} \end{bmatrix} = BA$$

$(d \times r) (r \times k)$

$d \leftarrow 3 \quad k \leftarrow 3$

$$= \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix} \cdot (a_1 \ a_2 \ a_3)$$

$(3 \times 1) \quad (1 \times 3)$

or

$$= \begin{pmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \\ b_{31} & b_{32} \end{pmatrix} \begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \end{pmatrix}$$

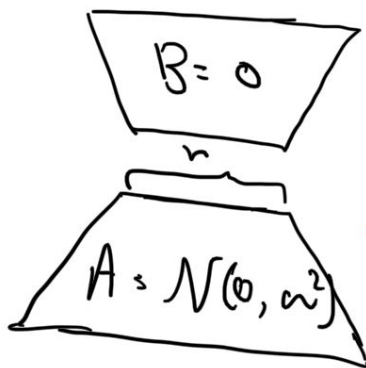
$(3 \times 2) \quad (2 \times 3)$

LoRA Tuning

$$\Delta W = \begin{bmatrix} \Delta w_{11} & \Delta w_{12} & \Delta w_{13} \\ \Delta w_{21} & \Delta w_{22} & \Delta w_{23} \\ \Delta w_{31} & \Delta w_{32} & \Delta w_{33} \end{bmatrix} = B A$$

$(d \times r) (r \times k)$

$3 \leftarrow d \times k \rightarrow 3$



$B = 0$

$A = \mathcal{N}(0, \sigma^2)$

or

$$= \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix} \cdot (a_1 \ a_2 \ a_3)$$


(3×1) (1×3)

$$= \begin{pmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \\ b_{31} & b_{32} \end{pmatrix} \begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \end{pmatrix}$$

(3×2) (2×3)

LoRA Tuning


During training



$$h = Wx + BAx$$

$$h = \underbrace{(W + BA)}_{W_{merged}}x$$

After training





Quiz

LoRA Tuning


- A) Improves Training Memory footprint
- B) Improves Inference Memory footprint
- C) A & B

Quiz

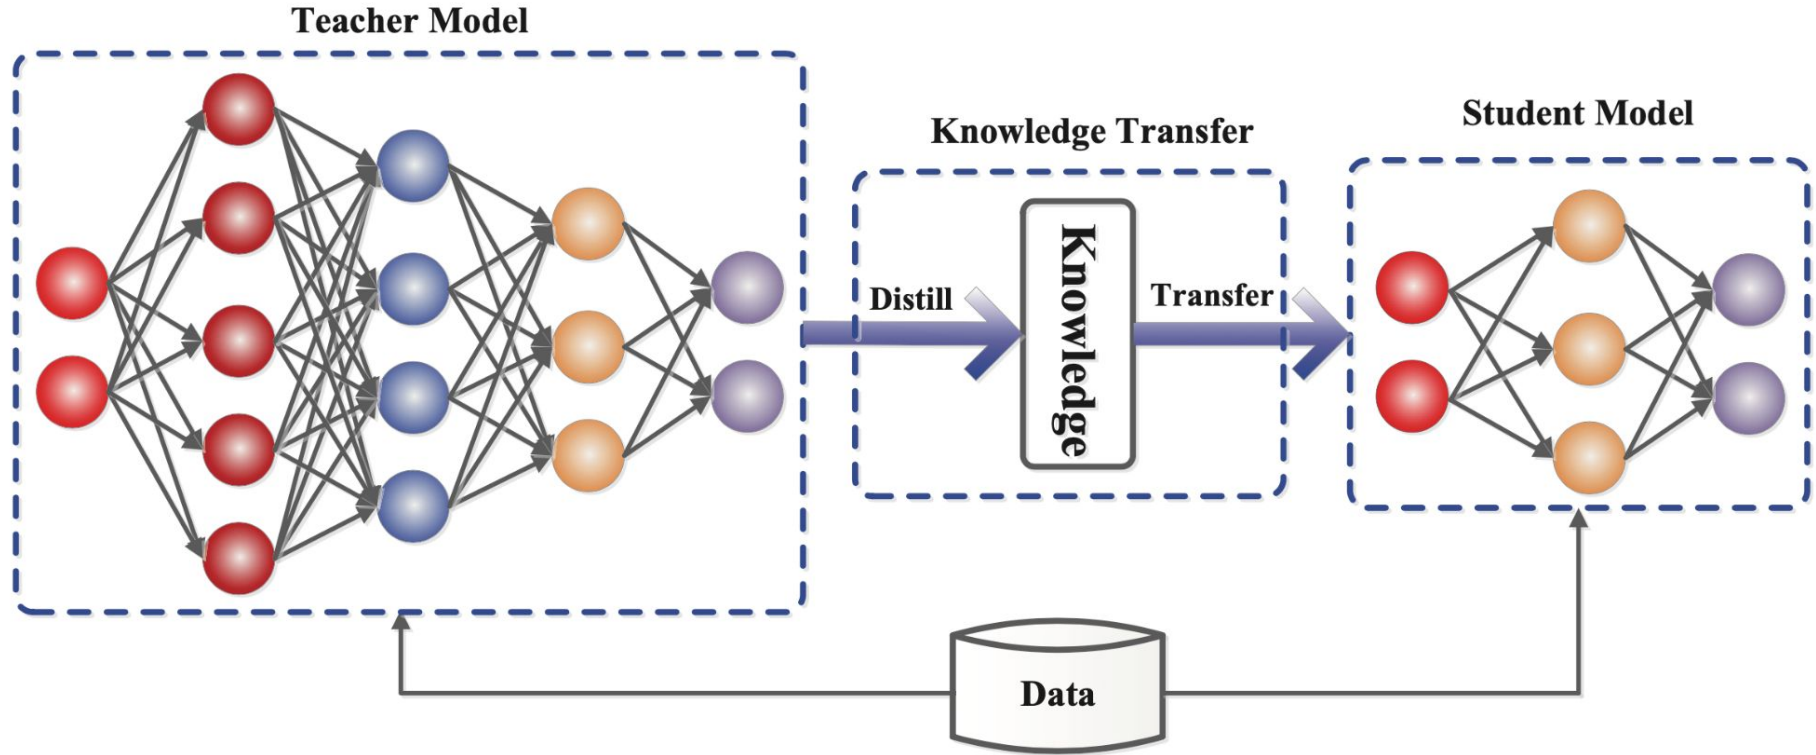
LoRA Tuning

- A) Improves Training Memory footprint
- B) Improves Inference Memory footprint
- C) A & B

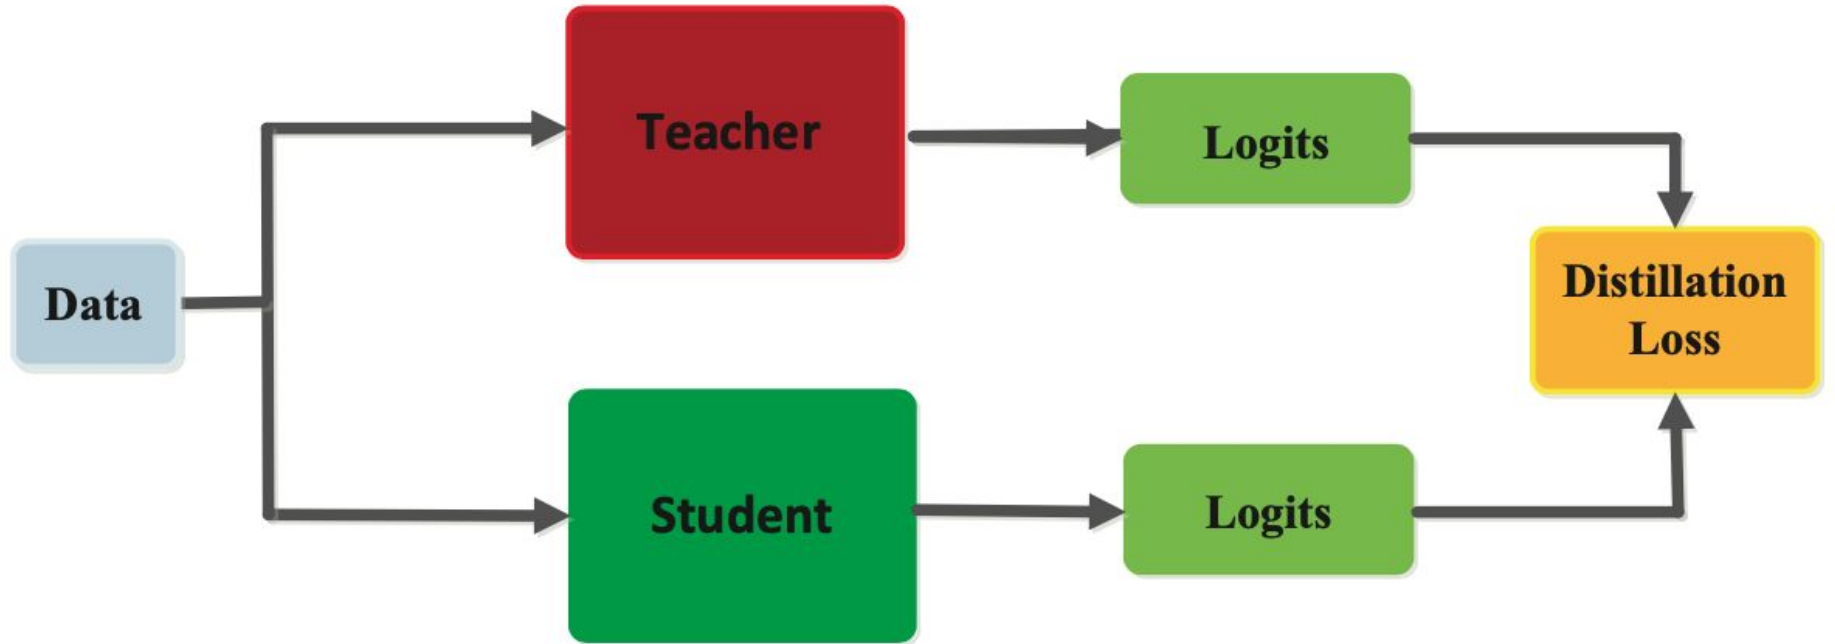
Knowledge Distillation



Knowledge Distillation



Response-based knowledge




Distillation

Honda SUV chair
0.6 0.35 0.05



...




i/p img 1

Honda SUV chair
0.39 0.59 0.02



...




i/p img 2

Distillation

student learner player
0.6 0.35 0.05



...




i/p txt 1

student learner player
0.39 0.59 0.02

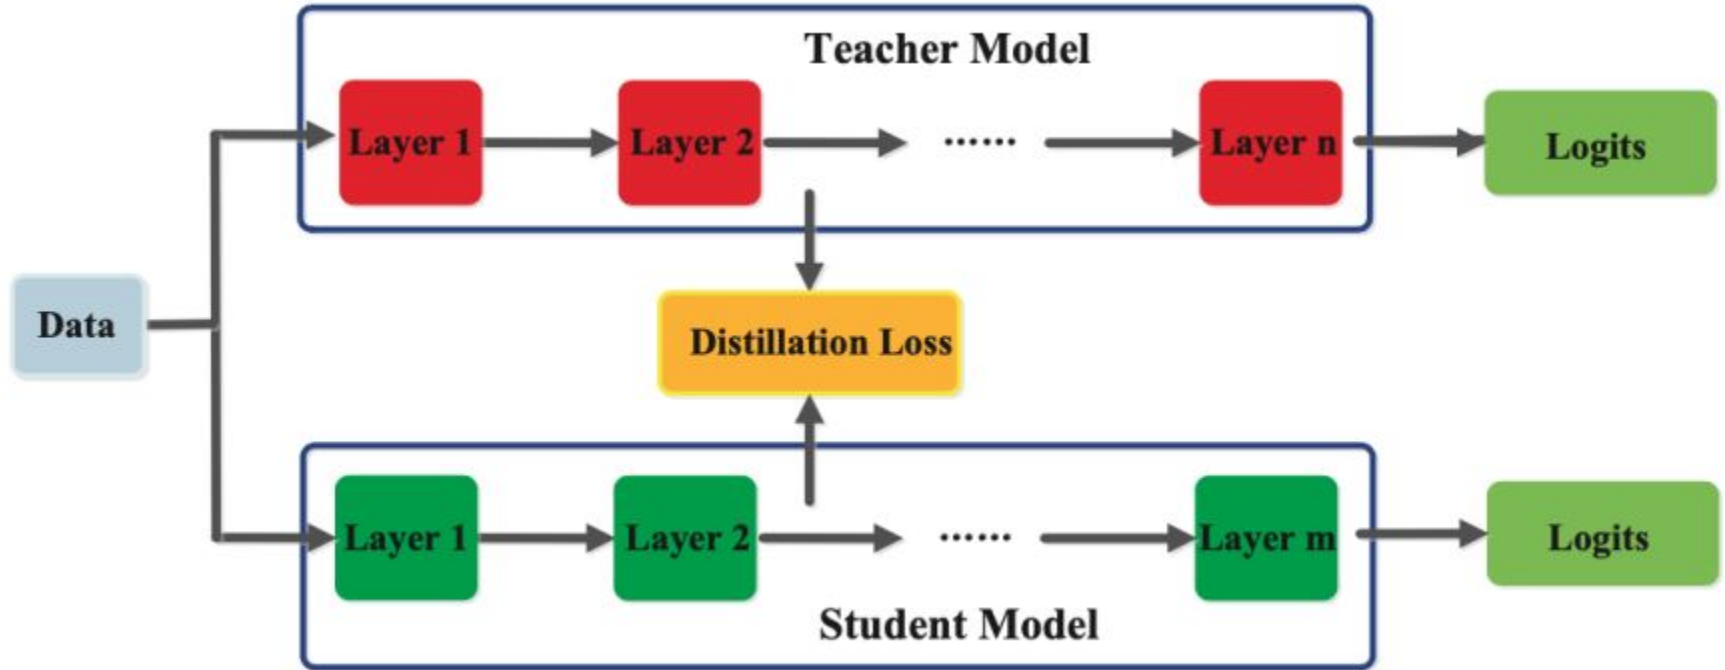


...



i/p txt 2

Feature-based knowledge



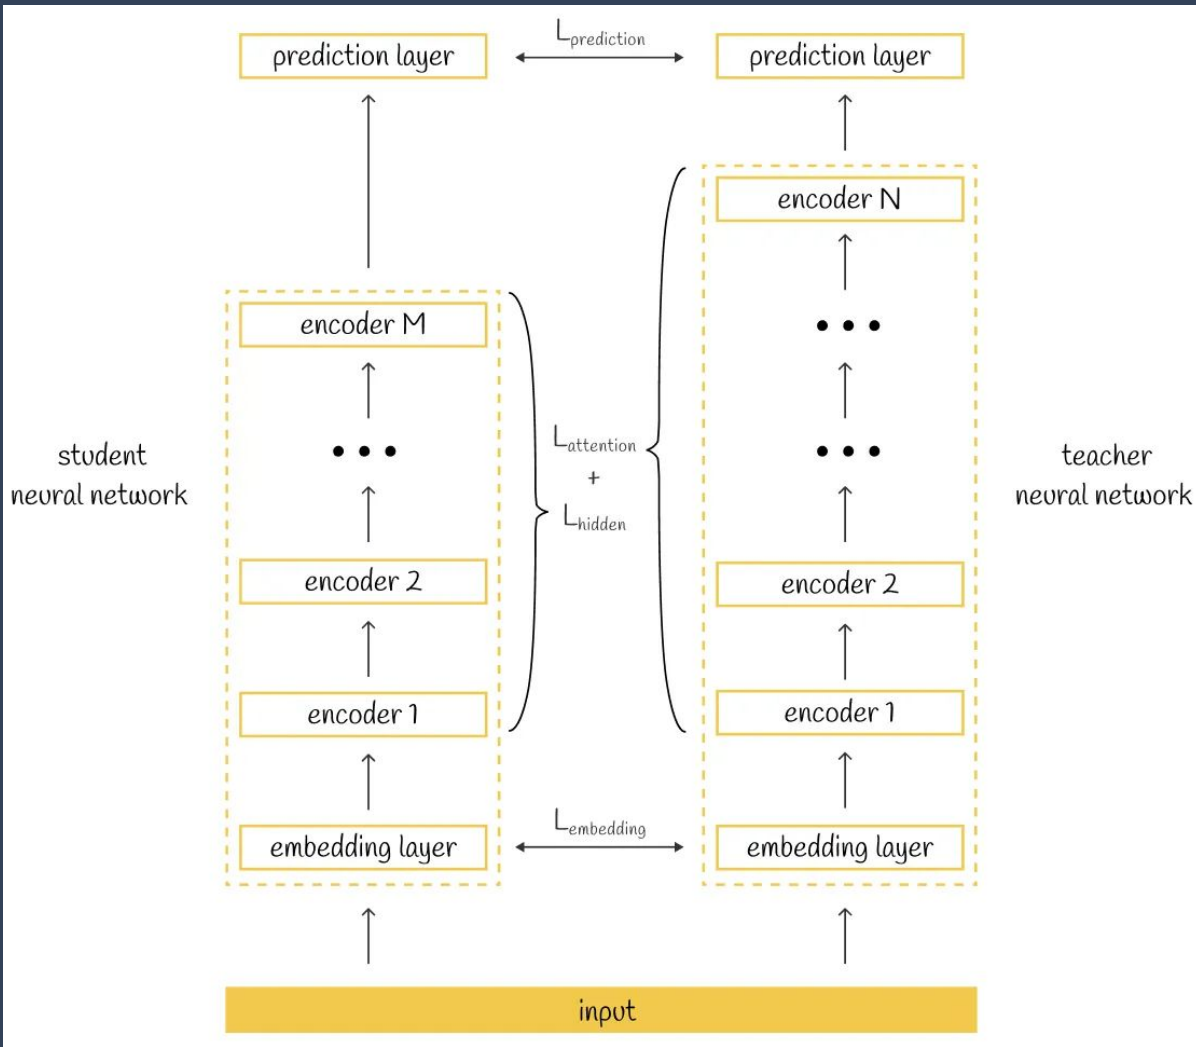
Demo Time

<https://colab.research.google.com/drive/1MH2kIPtcGzu0g2VkQ8lcCHFVeR6nAls7>

Distillation in LLM

- DistilBERT
 - The student model imitate the teacher model
 - Has about half the total number of parameters of BERT base and retains 95% of BERT's performances on the language understanding benchmark

TinyBert



Distillation: Practical Considerations

- The student model will be good at the distillation task only
- You still need unlabeled data
 - In some enterprise setups you may not be allowed to use customer data
- Model refresh



Quiz

Distillation Improves:

- A) Training latency
- B) Inference latency
- C) A & B

Quiz

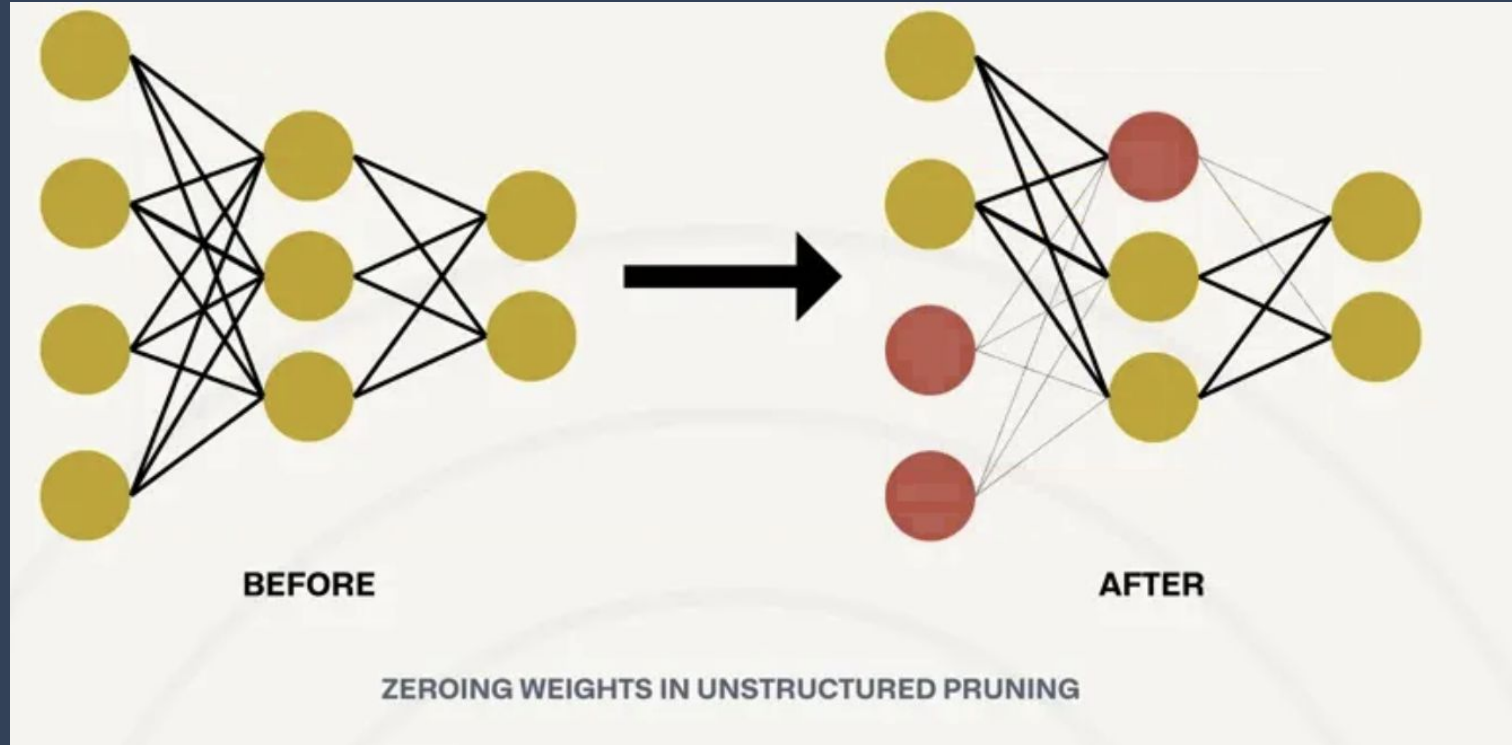
Distillation Improves:

- A) Training latency
- B) Inference latency
- C) A & B


Model Pruning

- The reasoning is that:
“Neural networks have an excess of parameters needed to generalize well and make accurate predictions”
- Then, we should drop these extra parameters









Model Pruning: Unstructured



Model Pruning: Structured





Model Pruning

- At training time
 -   More efficient as the model is trained with sparsity at training time
 -   Can be more complex
- Post training time
 -   Simpler to implement and adjust per use-case
 -   May require fine-tuning

Model Quantization

Quantization





Quiz

Model Pruning and Quantization Improve

- A) Training latency
- B) Inference latency
- C) A & B

Quiz

Model Pruning and Quantization Improve

- A) Training latency
- B) Inference latency**
- C) A & B

Is that it?

No, there are many many other optimizations

- Speculative decoding
- Linear attention
- Model parallelism
- ...

Summary and recap

- We went through the decoder details
- Sampling strategies
- Optimizations
 - KV Caching
 - LoRA
 - Distillation
 - Pruning
 - Quantizations

Questions & Discussions

References

- https://dugas.ch/artificial_curiosity/GPT_architecture.html
- <https://towardsdatascience.com/large-language-models-gpt-1-generative-pre-trained-transformer-7b895f296d3b>
- <https://medium.com/@yumo-bai/why-are-most-llms-decoder-only-590c903e4789>
- <https://medium.com/@javaid.nabi/all-you-need-to-know-about-llm-text-generation-03b138e0ed19>
- https://mlabonne.github.io/blog/posts/2023-06-07-Decoding_strategies.html
- <https://www.appypie.com/blog/hardware-requirements-for-llm-training#Hardware>
- <https://sujayskumar.com/2020/04/24/reducing-model-size/>
- <https://snorkel.ai/llm-distillation-demystified-a-complete-guide/>
- <https://neptune.ai/blog/knowledge-distillation>
- <https://medium.com/huggingface/distilbert-8cf3380435b5>
- <https://www.datature.io/blog/a-comprehensive-guide-to-neural-network-model-pruning>
- <https://www.youtube.com/watch?v=80bIUggRIf4>
- <https://medium.com/@joaolages/kv-caching-explained-276520203249>
- https://huggingface.co/docs/peft/main/en/conceptual_guides/lora
- <https://www.youtube.com/watch?v=KEv-F5UkhxU>
- <https://www.youtube.com/watch?v=KEv-F5UkhxU>